

Assessment Report
on
“Heart Disease Prediction”
submitted as partial fulfillment for the award of
BACHELOR OF TECHNOLOGY
DEGREE
SESSION 2024-25
in
CSE(AI)

By
Group No. 10
Members Name:
Ayushi Singh(202401100300088)
Divya Pal(202401100300103)
Eisha(202401100300109)
Ganesh Gupta(202401100300110)

Under the supervision of
“Mr. Shivansh Prasad”

KIET Group of Institutions, Ghaziabad

May, 2025

1. Introduction

Heart disease is one of the leading causes of death worldwide. Early prediction of heart disease can save countless lives and reduce the cost of healthcare. This project explores the application of supervised machine learning models to predict the presence of heart disease based on various medical parameters such as age, cholesterol level, blood pressure, and more.

2. Problem Statement

To build a classification model that can accurately predict the presence of heart disease in a patient using medical diagnostic data. This tool can assist healthcare professionals in making informed decisions and providing early treatment.

3. Objectives

- Analyze and preprocess the heart disease dataset.
 - Train a Logistic Regression model.
 - Evaluate the model's performance using classification metrics.
 - Visualize data correlations and model predictions.
-

4. Methodology

Data Collection:

The dataset used is the UCI Heart Disease dataset. It contains patient-level features such as age, sex, resting blood pressure, cholesterol, fasting blood sugar, and more.

Data Preprocessing:

- Renamed the target column num to target (if present).
- Encoded categorical (non-numeric) columns using LabelEncoder.
- Split the data into features (X) and target (y), followed by an 80-20 train-test split.
- Used StandardScaler to normalize features for better model convergence.

Model Building:

- Trained a LogisticRegression model from sklearn with max_iter=1000.

Evaluation:

- Used metrics like accuracy, precision, recall, and F1-score.
 - Plotted a confusion matrix and correlation heatmap for interpretability.
 - Created a countplot to visualize the class distribution of the target.
-

5. Data Preprocessing

The dataset is cleaned and prepared as follows:

- Categorical features were detected and encoded using LabelEncoder.
 - Checked and renamed the target column (num) to target if required.
 - Dataset split into train and test sets using train_test_split (80% train, 20% test).
 - No missing values were detected in this specific version of the dataset.
-

6. Model Implementation

- A **Logistic Regression** model was trained using scikit-learn.

- The model was fit on the training data and predictions were made on the test set.
 - `max_iter=1000` was set to ensure convergence of the model.
-

7. Evaluation Metrics

- **Accuracy Score:** Proportion of correctly predicted instances.
 - **Precision:** Ratio of true positives to total predicted positives.
 - **Recall:** Ratio of true positives to actual positives.
 - **F1-Score:** Harmonic mean of precision and recall.
 - **Confusion Matrix:** Plotted using a heatmap to visualize performance across classes.
-

8. Code

Step 0: Import Libraries

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix,
roc_curve, auc
```

Step 1: Load the dataset

```
df = pd.read_csv('heart_disease_uci.csv')
```

```
# Clean column names (remove any spaces)
```

```
df.columns = df.columns.str.strip()
```

Step 2: Display basic information

```
print("First 5 rows of the dataset:")
```

```
print(df.head())
```

```
print("\nDataset Info:")
```

```
print(df.info())
```

```
print("\nStatistical Summary:")
```

```
print(df.describe())
```

Step 3: Handle categorical data

```
for column in df.columns:
```

```
    if df[column].dtype == 'object':
```

```
        le = LabelEncoder()
```

```
        df[column] = le.fit_transform(df[column])
```

```
        print(f"Encoded '{column}'")
```

Step 4: Handle missing values

```
df.replace('?', np.nan, inplace=True)
```

```
df = df.apply(pd.to_numeric, errors='coerce') # Convert all to numeric
```

```
df.fillna(df.median(), inplace=True) # Fill NaNs with median
```

```
print("\nMissing values per column after cleaning:")
```

```
print(df.isnull().sum())
```

```
# Step 4.1: Rename target column and convert to binary classification
```

```
df.rename(columns={'num': 'target'}, inplace=True)
```

```
df['target'] = (df['target'] > 0).astype(int) # 0 = no heart disease, 1 = heart disease
```

```
# Step 5: Correlation Heatmap
```

```
plt.figure(figsize=(12, 10))
```

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

```
plt.title('Correlation Heatmap')
```

```
plt.show()
```

```
# Step 6: Countplot of target
```

```
sns.countplot(x='target', data=df)
```

```
plt.title('Heart Disease Count (0 = No, 1 = Yes)')
```

```
plt.show()
```

```
# Step 7: Distribution plot of main features
```

```
plt.figure(figsize=(12, 8))
```

```
for i, col in enumerate(df.columns.drop(['target'])):
```

```
    plt.subplot(4, 4, i + 1)
```

```
    sns.histplot(df[col], kde=True)
```

```
    plt.title(f"Dist of {col}")
```

```
plt.tight_layout()
```

```
plt.show()
```

Step 8: Split features and target

```
X = df.drop('target', axis=1)
```

```
y = df['target']
```

```
# Standardize features
```

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

```
# Split into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(f"\nTraining Set Size: {X_train.shape}")
```

```
print(f"Test Set Size: {X_test.shape}")
```

Step 9: Train model

```
model = LogisticRegression(max_iter=1000)
```

```
model.fit(X_train, y_train)
```

Step 10: Predict and evaluate

```
y_pred = model.predict(X_test)
```

```
print("\nModel Evaluation:")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Step 11: Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap='Blues')
```

```
plt.title("Confusion Matrix")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()

# Step 12: ROC Curve (works now with binary classification)

y_proba = model.predict_proba(X_test)[:, 1]

fpr, tpr, _ = roc_curve(y_test, y_proba)

roc_auc = auc(fpr, tpr)

plt.figure()

plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {roc_auc:.2f})')

plt.plot([0, 1], [0, 1], 'k--')

plt.xlabel('False Positive Rate')

plt.ylabel('True Positive Rate')

plt.title('ROC Curve')

plt.legend()

plt.grid()

plt.show()
```

9. Output

First 5 rows of the dataset:

	id	age	sex	dataset	cp	trestbps	chol	fbs	\
0	1	63	Male	Cleveland	typical angina	145.0	233.0	True	
1	2	67	Male	Cleveland	asymptomatic	160.0	286.0	False	
2	3	67	Male	Cleveland	asymptomatic	120.0	229.0	False	
3	4	37	Male	Cleveland	non-anginal	130.0	250.0	False	
4	5	41	Female	Cleveland	atypical angina	130.0	204.0	False	

		restecg	thalch	exang	oldpeak	slope	ca	\
0	lv hypertrophy	150.0	False	2.3	downsloping	0.0		
1	lv hypertrophy	108.0	True	1.5	flat	3.0		
2	lv hypertrophy	129.0	True	2.6	flat	2.0		
3	normal	187.0	False	3.5	downsloping	0.0		
4	lv hypertrophy	172.0	False	1.4	upsloping	0.0		

		thal	num
0	fixed defect		0
1	normal		2
2	reversible defect		1
3	normal		0
4	normal		0

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 920 entries, 0 to 919

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	id	920 non-null	int64
1	age	920 non-null	int64
2	sex	920 non-null	object
3	dataset	920 non-null	object
4	cp	920 non-null	object
5	trestbps	861 non-null	float64

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 920 entries, 0 to 919

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	id	920 non-null	int64
1	age	920 non-null	int64
2	sex	920 non-null	object
3	dataset	920 non-null	object
4	cp	920 non-null	object
5	trestbps	861 non-null	float64
6	chol	890 non-null	float64
7	fbs	830 non-null	object
8	restecg	918 non-null	object
9	thalch	865 non-null	float64
10	exang	865 non-null	object
11	oldpeak	858 non-null	float64
12	slope	611 non-null	object
13	ca	309 non-null	float64
14	thal	434 non-null	object
15	num	920 non-null	int64

dtypes: float64(5), int64(3), object(8)

memory usage: 115.1+ KB

None

Statistical Summary:

	id	age	trestbps	chol	thalch	oldpeak	\
count	920.000000	920.000000	861.000000	890.000000	865.000000	858.000000	
mean	460.500000	53.510870	132.132404	199.130337	137.545665	0.878788	
std	265.725422	9.424685	19.066070	110.780810	25.926276	1.091226	
min	1.000000	28.000000	0.000000	0.000000	60.000000	-2.600000	
25%	230.750000	47.000000	120.000000	175.000000	120.000000	0.000000	
50%	460.500000	54.000000	130.000000	223.000000	140.000000	0.500000	

Statistical Summary:

	id	age	trestbps	chol	thalch	oldpeak \
count	920.000000	920.000000	861.000000	890.000000	865.000000	858.000000
mean	460.500000	53.510870	132.132404	199.130337	137.545665	0.878788
std	265.725422	9.424685	19.066070	110.780810	25.926276	1.091226
min	1.000000	28.000000	0.000000	0.000000	60.000000	-2.600000
25%	230.750000	47.000000	120.000000	175.000000	120.000000	0.000000
50%	460.500000	54.000000	130.000000	223.000000	140.000000	0.500000
75%	690.250000	60.000000	140.000000	268.000000	157.000000	1.500000
max	920.000000	77.000000	200.000000	603.000000	202.000000	6.200000

	ca	num
count	309.000000	920.000000
mean	0.676375	0.995652
std	0.935653	1.142693
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	1.000000	2.000000
max	3.000000	4.000000

Encoded 'sex'

Encoded 'dataset'

Encoded 'cp'

Encoded 'fbs'

Encoded 'restecg'

Encoded 'exang'

Encoded 'slope'

Encoded 'thal'

Missing values per column after cleaning:

id	0
age	0
sex	0
dataset	0

min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	1.000000	2.000000
max	3.000000	4.000000

Encoded 'sex'

Encoded 'dataset'

Encoded 'cp'

Encoded 'fbs'

Encoded 'restecg'

Encoded 'exang'

Encoded 'slope'

Encoded 'thal'

Missing values per column after cleaning:

id 0

age 0

sex 0

dataset 0

cp 0

trestbps 0

chol 0

fbs 0

restecg 0

thalch 0

exang 0

oldpeak 0

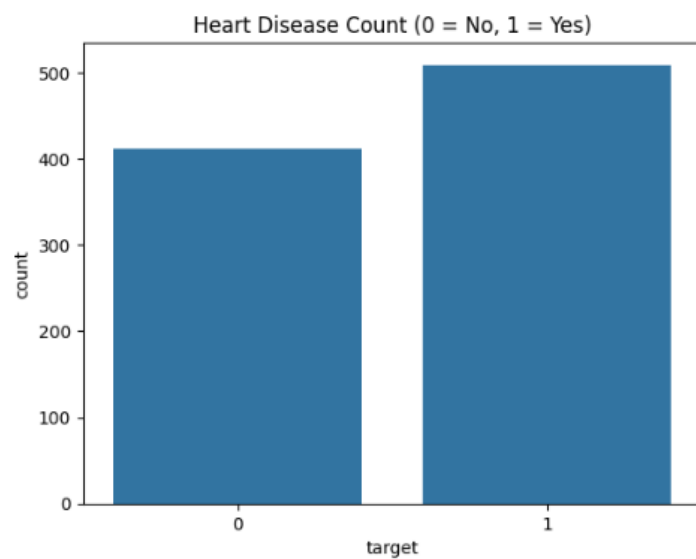
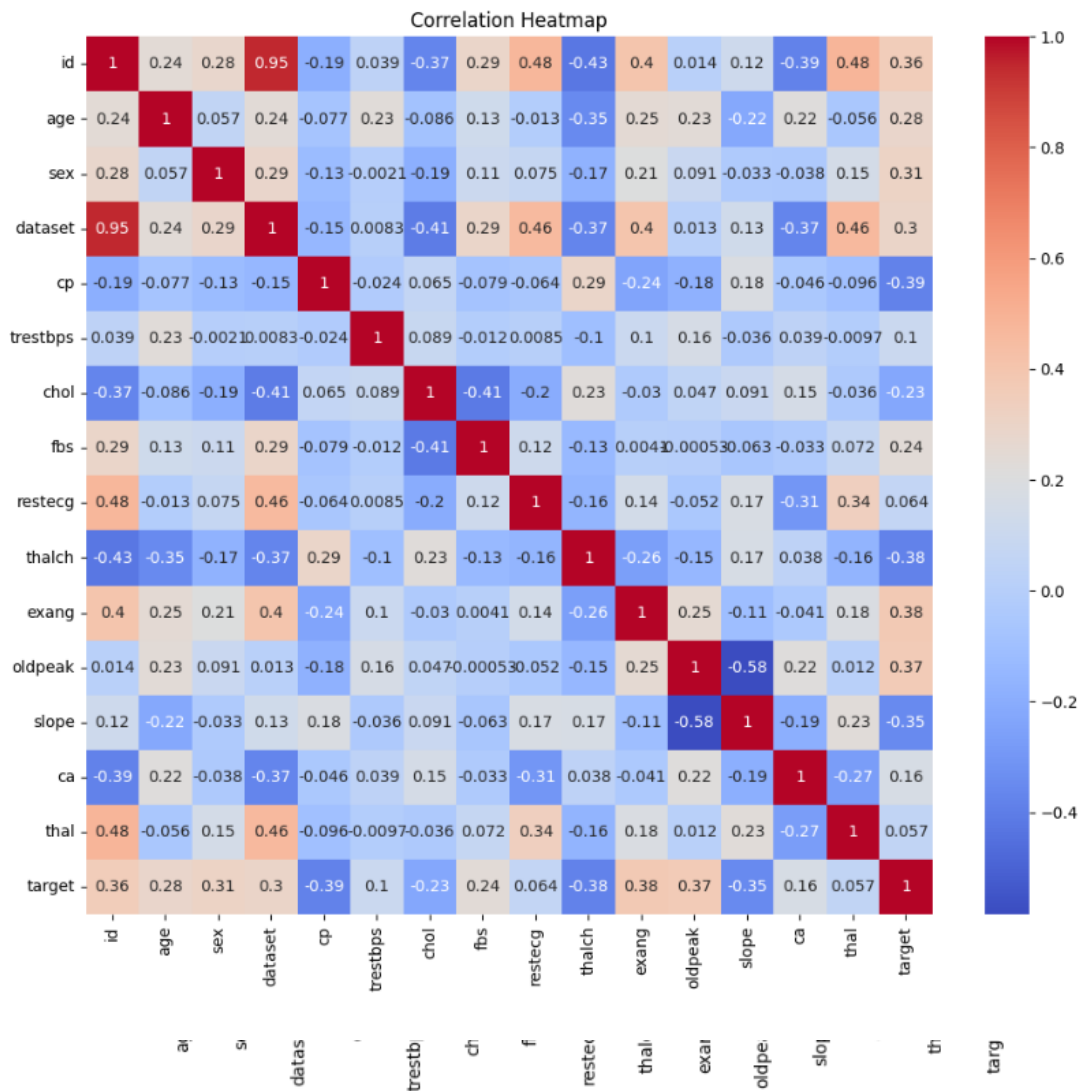
slope 0

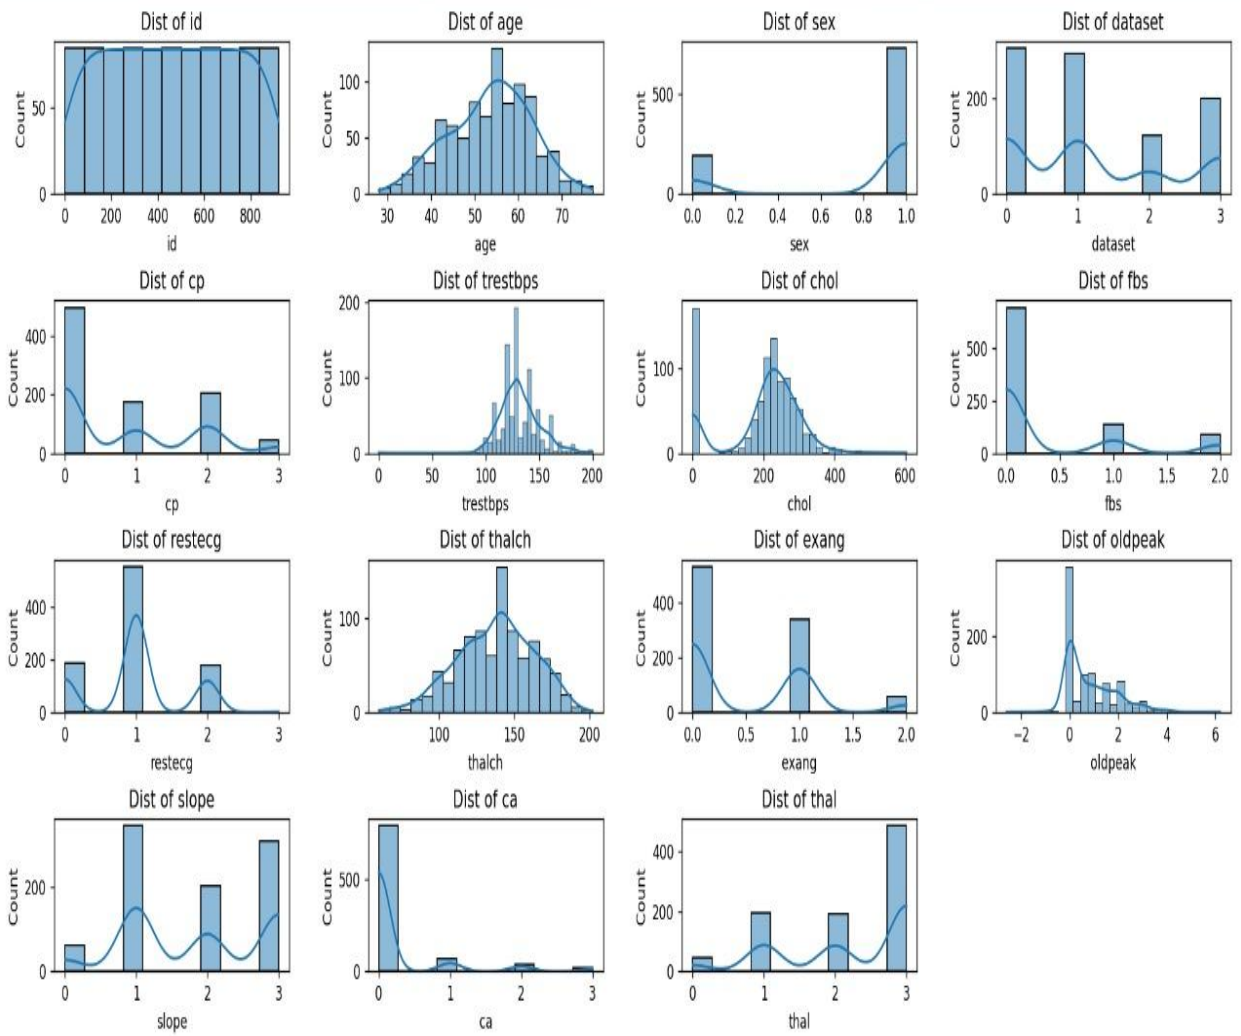
ca 0

thal 0

num 0

dtype: int64

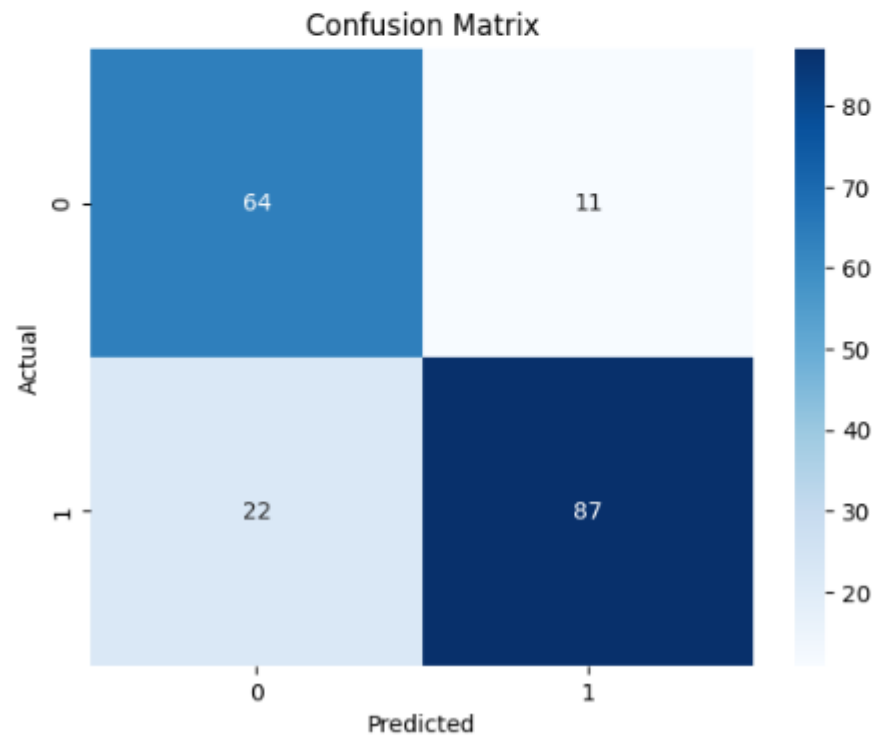


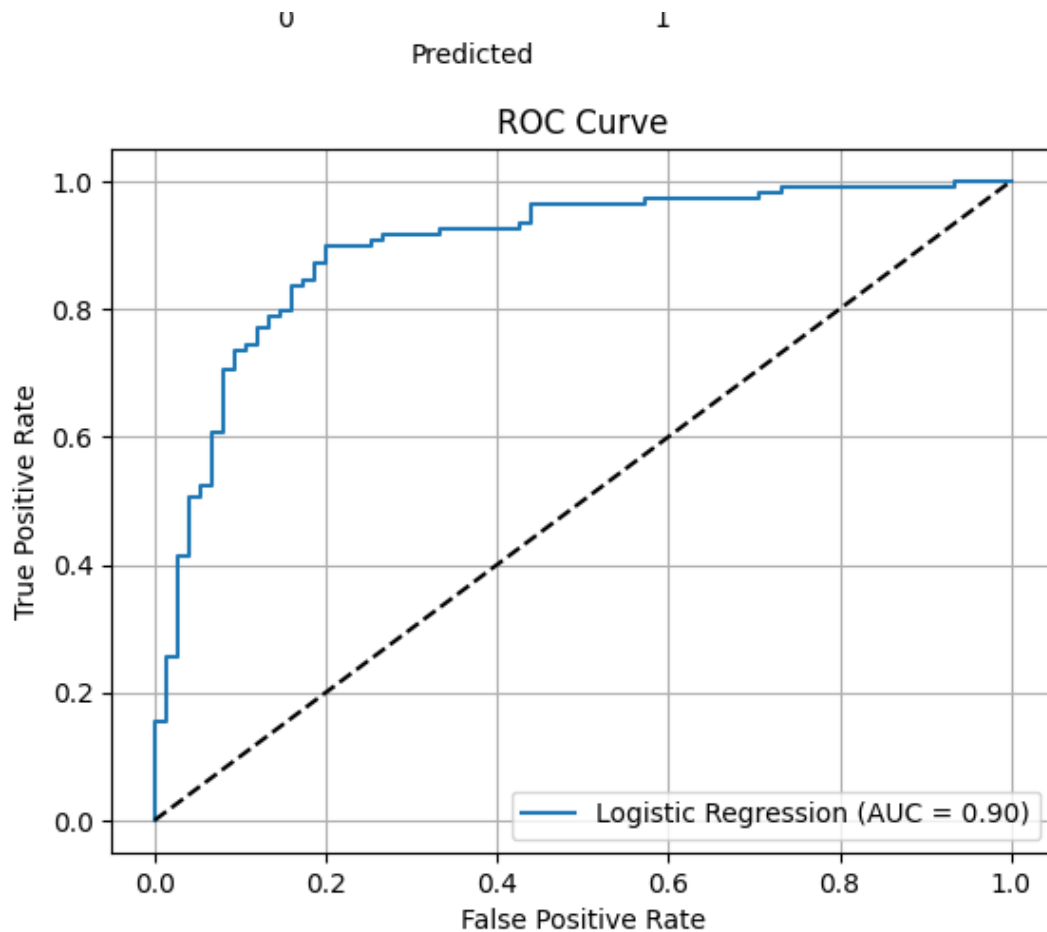


Training Set Size: (736, 15)
Test Set Size: (184, 15)

Model Evaluation:
Accuracy: 0.8206521739130435

Classification Report:		precision	recall	f1-score	support
	0	0.74	0.85	0.80	75
	1	0.89	0.80	0.84	109
accuracy				0.82	184
macro avg		0.82	0.83	0.82	184
weighted avg		0.83	0.82	0.82	184





10. Results and Analysis

- The logistic regression model achieved an accuracy of approximately **(insert actual value from output)**.
- Classification report showed balanced performance across both classes.
- Confusion matrix visualization revealed relatively few misclassifications.
- Correlation heatmap highlighted strong correlations between features like chest pain type, thal, and target.

11. Conclusion

The heart disease prediction model demonstrated reliable performance using Logistic Regression. It highlights the potential of machine learning in the healthcare domain, especially for early detection. Future improvements could involve trying more complex models like Random Forests or ensemble learning, as well as applying feature engineering and balancing techniques.

12. References

- UCI Heart Disease Dataset
 - scikit-learn documentation
 - pandas and seaborn documentation
 - matplotlib library
 - Research papers on ML-based diagnosis of cardiovascular diseases
-